

Model and estimators for partial least squares.

Running title: Model and estimators for PLS.

Helland, I.S., University of Oslo, Sæbø, S., Almøy, T. and Rimal, R.,
Norwegian University of Life Sciences

March 22, 2017

Abstract

Partial least squares regression has been a very popular method for prediction, first used by chemometricians, but also now by statisticians in a number of applied fields. The method can in a natural way be connected to a statistical model which now has been extended and further developed in terms of an envelope model. Concentrating on the univariate case, the model is here approached from several points of view. Several estimators of the regression vector in this model are defined, including the ordinary PLS estimator, the maximum likelihood envelope estimator and a recently proposed Bayes PLS estimator. The different estimators are compared with respect to prediction error by systematic simulations. The main conclusion from the simulations seems to be that Bayes PLS performs well compared to the other methods in virtually all cases.

Key words: Bayes PLS estimator; envelope model; maximum likelihood envelope estimator; partial least squares; partial least squares model; prediction; simulation.

1 Introduction

Supervised learning from multivariate data is a central problem area in applied statistics. Specifically, let our task be to predict a single variable y from a p -dimensional variable x , having data on n units. A large number of learning methods are discussed in Hastie et al. (2009), mainly under the ordinary multiple regression model. It is part of the philosophy of the present paper that for near collinear data, a *reduction* of the regression model is called for.

Here we will take as our point of departure a well defined reduction of the random x regression model, the envelope model of Cook et al. (2013). In op. cit. it has been shown in

detail that the envelope model in the case of reduction in the x -variable is equivalent to the partial least square population model of Helland (1990), and that a corresponding result is also valid in the case of a multivariate y .

Partial least squares (PLS) regression has had a vigorous development in the chemometric literature since it was proposed by Herman and Svante Wold and by Harald Martens (Wold et al., 1984; Martens and Naes, 1992). The method has been extended in several directions and its applications have been expanded to an increasing number of fields, for instance genomic data (Boulesteix and Strimmer, 2007). Both these issues have been discussed in detail in a recent paper by Mehmood and Ahmed (2016), where a wealth of further references may be found.

In the beginning the PLS method was to some extent neglected or turned down by statisticians (an exception among others was Frank and Friedman, 1993), but it is now included as a tool among other biased regression methods by applied statisticians. For a general discussion paper with contributions both from mathematical statisticians and chemometricians, see Sundberg (1999).

Indeed there was a difference in culture between chemometricians and statisticians then, and this difference still exists to a large extent. A statement by Munck et al. (2010) illustrates this, as seen from one side: 'If chemometrics in its historical development had been limited to follow current scientific (and statistical) theories there would have been minimal progress in its wide applications today.'

This difference in culture may be related to the concepts of creativity and rigor, qualities which to some extent may be called complementary. One could say that one culture puts more emphasis on creativity, the other on rigor. Of course, this is a huge simplification. First, there is a lot of creativity among statisticians, also mathematical statisticians. Secondly, one should emphasize that precise thinking also should influence practice. A case of point is the following: Chung and Keles (2010) recently proved that the PLS regression vector is inconsistent when $p/n \rightarrow k > 0$ under a wide set of conditions. This result is probably not too well known among chemometricians; some may have a tendency to put too much confidence in PLS regression when $p \sim n$ or $p > n$. It is to be emphasized that the inconsistency result of Chung and Keles (2010) is only concerned with estimation of the regression vector. The mathematical properties of PLS as a *prediction* method when $p > n$ are largely unknown, as far as we know, although there is much positive empirical evidence among applied researchers on these properties.

It is true that chemometricians have had a leading edge in the development of PLS and of certain multivariate methods, in particular with respect to visualization etc., and they still may be ahead of us in this sense. (However, there is a growing literature in statistical journals on the envelope model and on estimation in this model; see below.)

Accepting this, an important general question is what mathematical statisticians can contribute with in this development. A vital aspect in the history of statistics is the interplay between model and estimators. Once a model is formulated, one can in principle think of several estimators in this model. Thus in the PLS/envelope model one can certainly discuss other estimators than the usual PLS regression estimator, which can be seen as originating by replacing population (co-)variances in the model by sample (co-)variances. Two examples are the maximum likelihood estimator of Cook et al. (2013), see also Cook et al. (2015), Cook et al. (2016) and Cook and Zhang (2016); and the Bayesian estimator of Helland et al. (2012). By simulation, both these estimators have performed well compared to PLS regression, but they have their disadvantages. The maximum likelihood estimator can not be used in the case when the data matrix has rank less than p , and the Bayesian estimator requires heavy computations, in particular when p is large.

To compare estimators we make vital use of the recently developed simulation package `simrel`; see Sæbø et al. (2015). A main purpose of the present paper is to illustrate the interplay between mathematical arguments on the one hand and systematic simulations on the other hand in an area where it is difficult to obtain results by purely analytical means.

It is important to emphasize that this paper is based upon reduction of the *random x* regression model. When considering latent variables from PLS, and when considering near collinearity in the observed x -variables, it is natural to treat these x -variables as random. It is our philosophy that this is also the best way to look upon model reduction. On the other hand, in the context of prediction, one could argue that one should condition upon the x -variables and consider them as fixed. A prominent paper on PLS regression, taking fixed x -variables in the basic model, is Krämer and Sugiyama (2012), where further references can be found.

In recent years there has been a rapidly growing literature on the envelope model. In addition to the maximum likelihood estimation paper mentioned above, the most important papers seem to be Cook and Zhang (2015b), where simultaneous reduction in the x -space and y -space is proposed, and Cook and Zhang (2015a), where extensions to other regression methods than linear regression are discussed. More references can be found in these papers.

Model reduction in regression models is discussed in general from the point of view of rotations in the x -space in Helland (2001) and from a different point of view in Helland (2000).

The plan of this paper is as follows: In Section 2 we formulate the model in 4 different ways which can be shown to be equivalent. In Section 3 we define 4 different estimators in the model. In Section 4 we ask the question if the ordinary PLS estimator with m components can be improved by forcing the weight vector at step $m + 1$ to vanish; the answer turns out to be negative. Then in Section 5 we define the Bayes PLS estimator of Helland et al. (2012). In Section 6 we describe the simulations done for comparison of estimators with respect to prediction error, and in Section 7 we give the results of the simulations. Finally, Section 8 is a discussion section.

2 The model; several formulations

Take as a point of departure the linear model

$$y = \mu_y + \beta'(x - \mu_x) + \epsilon, \quad (1)$$

where β and x are p -dimensional, and where the random predictor x has mean μ_x and covariance matrix Σ_{xx} , for simplicity assumed nonsingular here. (This can be relaxed to assuming $\beta \in \text{span}(\Sigma_{xx})$ in the case where this matrix is singular; see Cook et al. (2013), and also C below.) Independently, ϵ is distributed with mean 0 and variance σ^2 . When doing prediction from this model for near collinear data, a model reduction may be called for. Throughout this paper, a definite m -dimensional model reduction, which may be formalized in several equivalent ways, will be used. When this model holds, we say that we have an envelope model of dimension m , or that there are m relevant components for prediction in the model.

- A. Given a subspace \mathcal{T} of R^p , let $P_{\mathcal{T}}$ be the projection upon \mathcal{T} , and let $Q_{\mathcal{T}}$ be the projection orthogonal to \mathcal{T} . For simplicity discuss the case where $\mu_x = 0$. Let now \mathcal{T} be the smallest space such that (i) $Q_{\mathcal{T}}x$ is uncorrelated with $P_{\mathcal{T}}x$; (ii) y is uncorrelated with $Q_{\mathcal{T}}x$ given $P_{\mathcal{T}}x$. In this case we may say that $Q_{\mathcal{T}}x$ contains no information about y , neither directly nor through $P_{\mathcal{T}}x$. Here \mathcal{T} is a subspace of the x -space. Let \mathcal{S} be the corresponding dual subspace of the parameter-space: Define \mathcal{S}^{\perp} as the set of γ such that $\gamma'x = 0$ for all $x \in \mathcal{T}$, and let \mathcal{S} be the space perpendicular to \mathcal{S}^{\perp} . Consider model reduction to $\beta \in \mathcal{S}$,

which is equivalent to $x \in \mathcal{T}$, since β and x only occur in the model in the combination $\beta'x$.

- B. Here is an algebraic characterization which turns out to be equivalent. For a matrix M define MS as the space of vectors Mz as z runs through S , and let S^\perp be the space perpendicular to S . Let now \mathcal{S} be the smallest space in R^p such that (i) both $\Sigma_{xx}\mathcal{S} \subseteq \mathcal{S}$ and $\Sigma_{xx}\mathcal{S}^\perp \subseteq \mathcal{S}^\perp$; (ii) $\text{span}(\beta) \subseteq \mathcal{S}$. In this case we say that \mathcal{S} is the envelope of $\text{span}(\beta)$. It can be shown (Cook et al., 2010) that the envelope always exists as the smallest space with the stated properties.
- C. The regression vector β can always be expanded in terms of the eigenvectors d_i of Σ_{xx} :

$$\beta = \sum_{i=1}^p \gamma_i d_i. \quad (2)$$

Assume that this sum can be reduced to exactly m non-zero terms: $\beta = \sum_{i=1}^m \gamma_i d_i$, where the d_i correspond to different eigenvalues of Σ_{xx} . We then say that there are m relevant components for prediction in the model. This reduction can be imagined to take place by two mechanisms: 1) Some of the γ_i 's are really zero. 2) There are coinciding eigenvalues in Σ_{xx} . Then one can rotate such that it is enough with one eigenvector for each eigenspace in the sum. In this approach it is important that we only know that there are m non-zero terms in the sum, not which terms that are non-zero. For a closer discussion of this, see Naes and Helland (1993) and Helland and Almøy (1994).

- D. Consider the population version of the well known PLS algorithm: Take $e_0 = x - \mu_x$, $f_0 = y - \mu_y$, and for $a = 1, 2, \dots, m$ compute successively:

$$w_a = \text{cov}(e_{a-1}, f_{a-1}), \quad t_a = w_a' e_{a-1}, \quad (3)$$

$$p_a = \text{cov}(e_{a-1}, t_a) / \text{var}(t_a), \quad q_a = \text{cov}(f_{a-1}, t_a) / \text{var}(t_a), \quad (4)$$

$$e_a = e_{a-1} - p_a t_a, \quad f_a = f_{a-1} - q_a t_a.$$

It can be proved (Helland, 1990), and is important in this connection that under the reduced model C, this algorithm stops automatically after m steps when $m < p$: It stops because $w_{m+1} = \text{cov}(e_m, f_m) = 0$. After those m steps we get the representations

$$\mathbf{x} = \boldsymbol{\mu}_x + \mathbf{p}_1 t_1 + \dots + \mathbf{p}_m t_m + \mathbf{e}_m, \quad y = \mu_y + q_1 t_1 + \dots + q_m t_m + f_m \quad (5)$$

with the corresponding PLS population prediction

$$y_{m,PLS} = \mu_y + q_1 t_1 + \dots + q_m t_m = \mu_y + \boldsymbol{\beta}'_{m,PLS}(\mathbf{x} - \boldsymbol{\mu}_x).$$

Theorem 1. (Cook et al., 2013; Helland, 1990)

- (a) *The two conditions A and B on the space \mathcal{S} are equivalent.*
- (b) *The models formulated by C and D are equivalent.*
- (c) *When there are m relevant components for prediction, the envelope space \mathcal{S} has dimension m , and \mathcal{S} can be taken as $\text{span}(\mathbf{w}_1, \dots, \mathbf{w}_m) = \text{span}(\mathbf{d}_1, \dots, \mathbf{d}_m)$.*
- (d) *When the envelope space has dimension m , there are m relevant components for prediction.*
- (e) *In this case we have $\boldsymbol{\beta}_{m,PLS} = \boldsymbol{\beta}$.*

In this sense all the model formulations A-D are equivalent; they describe the same reduced model. In Cook et al. (2013), a fifth equivalent formulation in terms of a Krylov sequence is also given. Being a reduced model that can be motivated in so many different ways, it is definitively of interest to find a good estimator of the regression vector $\boldsymbol{\beta}$ under this model.

3 Estimators in the PLS/envelope model

Now that the PLS model is introduced, we will start to look at estimators of the parameters in this model, in particular estimators of $\boldsymbol{\beta}$, which will give prediction. Of special interest is estimators that perform well in the case of near collinear data. Some estimators are already known from the literature.

- a. The ordinary PLS estimator can be introduced as follows: With data (\mathbf{X}, \mathbf{y}) , take initial values $\mathbf{E}_0 = \mathbf{X} - \bar{\mathbf{x}}\mathbf{1}'$ and $\mathbf{f}_0 = \mathbf{y} - \bar{y}\mathbf{1}$. Run the population PLS algorithm for A steps with population (co-)variances replaced by sample (co-)variances. Ordinarily A is found by cross-validation or by similar means. Note that from D in subsection 2.1, the m -step

PLS model is characterized by $\mathbf{w}_{m+1} = \text{cov}(\mathbf{e}_m, f_m) = \mathbf{0}$. Theoretically, when $A = m$, we can not expect the sample weights $\hat{\mathbf{w}}_{m+1}$ to be zero. However, since any continuous function of the sample covariances and variances is consistent for the same function of the population covariances and variances, and since $\hat{\mathbf{w}}_{m+1}$ through the PLS algorithm is such a function and since $\mathbf{w}_{m+1} = \mathbf{0}$, we will have $\lim_{n \rightarrow \infty} \hat{\mathbf{w}}_{m+1} = \mathbf{0}$ almost surely as $n \rightarrow \infty$.

- b. The sparse regression SPLS of Chun and Keleş (2010). This requires two effective tuning parameters, and it also aims at variable selection. SPLS seems to be better than ordinary PLS in certain cases, also when variable selection is not an issue.
- c. When $\mathbf{S} = (\mathbf{X} - \bar{\mathbf{x}}\mathbf{1}')'(\mathbf{X} - \bar{\mathbf{x}}\mathbf{1}')$ has rank p , which specifically requires $n > p$, the maximum likelihood estimator of β under the multinormal envelope model was given in Cook et al. (2013). This estimator is of course very useful, but it cannot be used for small n . Modifications of the maximum likelihood estimator which cover also this case, were recently indicated by Cook et al. (2015). That paper also gives a MATLAB toolbox for maximum likelihood estimation in the envelope model and in several generalizations of this model. A faster algorithm for maximum likelihood estimation is discussed in Cook and Zhang (2016); an even faster algorithm and an R-package was recently described by Cook et al. (2016).
- d. Under a specific rotation-invariant prior, the Bayes estimator of β under the model with m relevant components was given in Helland et al. (2012). This estimator was shown to be close to the best equivariant estimator, but it requires heavy computation.

By simulation both the maximum likelihood estimator c and the Bayes estimator d were shown to perform very well compared to the PLS estimator a. These two estimators require a multinormal distribution of the data. However, both the chemometric tradition and the envelope model of Cook et al. (2010, 2013) demand no detailed distributional assumptions.

4 Can a better estimator be found by simple means?

The m step PLS model is characterized by the constraint $\mathbf{w}_{m+1} = \text{cov}(\mathbf{e}_m, f_m) = \mathbf{0}$. However, in the sample PLS algorithm, $\hat{\mathbf{w}}_{m+1}$ is a continuous random variable if the data are continuous.

Hence almost surely $\hat{\mathbf{w}}_{m+1} \neq \mathbf{w}_{m+1} = \mathbf{0}$. This means that the estimator of the vector of PLS parameter falls outside the corresponding parameter space. On the other hand, by standard statistical theory, the maximum likelihood estimator and any Bayes estimator are always in the parameter space.

In this Section, we ask the question whether we can improve the PLS algorithm in some way such that $\hat{\mathbf{w}}_{m+1} = \mathbf{0}$ for the improved algorithm. That is, we seek modified weights $\hat{\mathbf{w}}_1, \dots, \hat{\mathbf{w}}_m$ such that $\hat{\mathbf{w}}_{m+1} = \mathbf{0}$ in the modified algorithm. Unfortunately the answer to this question is no. This programme is only possible when \mathbf{S} is invertible, and then it by necessity leads to the least squares solution. Let $\hat{\mathbf{W}}_A = (\hat{\mathbf{w}}_1, \dots, \hat{\mathbf{w}}_A)$ for any A .

First we need some properties of the ordinary PLS algorithm.

Proposition 1. *At each step the PLS weights satisfy*

$$\hat{\mathbf{w}}_{A+1} = \mathbf{s} - \mathbf{S}\hat{\mathbf{W}}_A(\hat{\mathbf{W}}_A'\mathbf{S}\hat{\mathbf{W}}_A)^{-1}\hat{\mathbf{W}}_A'\mathbf{s}, \quad (6)$$

and the A step regression vector is

$$\hat{\beta}_A = \hat{\mathbf{W}}_A(\hat{\mathbf{W}}_A'\mathbf{S}\hat{\mathbf{W}}_A)^{-1}\hat{\mathbf{W}}_A'\mathbf{s}. \quad (7)$$

Proof. These relations were proved in Helland (1988) and were also used in Cook et al. (2013). □

Now fix m . To find an algorithm such that $\hat{\mathbf{w}}_{m+1} = \mathbf{0}$, we will have to modify the weights $\hat{\mathbf{w}}_1, \dots, \hat{\mathbf{w}}_m$. By definition we will call a restricted PLS (RPLS) prediction any method based on an estimator of β of the form (7) for $A = m$ with the proviso that 1) $\hat{\mathbf{W}}_m$ is modified in some way. 2) (6) holds for $A = m$, giving $\hat{\mathbf{w}}_{m+1} = \mathbf{0}$.

Theorem 2. *RPLS exists if and only if \mathbf{S} is invertible and $\mathbf{S}^{-1}\mathbf{s} \in \text{span}\hat{\mathbf{W}}_m$. In that case $\hat{\beta}$ is equal to the least squares estimator $\mathbf{S}^{-1}\mathbf{s}$.*

Proof. Assume that (6) holds for $A = m$ and $\hat{\mathbf{w}}_{m+1} = \mathbf{0}$. Then $\mathbf{s} = \mathbf{S}\hat{\mathbf{W}}_m(\hat{\mathbf{W}}_m'\mathbf{S}\hat{\mathbf{W}}_m)^{-1}\hat{\mathbf{W}}_m'\mathbf{s}$. This is possible for general \mathbf{s} only if \mathbf{S} is nonsingular, and then it is equivalent to $\mathbf{R}\sqrt{\mathbf{S}}^{-1}\mathbf{s} = \sqrt{\mathbf{S}}^{-1}\mathbf{s}$ with $\mathbf{R} = \mathbf{A}(\mathbf{A}'\mathbf{A})^{-1}\mathbf{A}'$, where $\mathbf{A} = \sqrt{\mathbf{S}}\hat{\mathbf{W}}_m$. Since \mathbf{R} is the projector upon $\text{span}(\mathbf{A})$, this is again equivalent to $\sqrt{\mathbf{S}}^{-1}\mathbf{s} \in \text{span}(\sqrt{\mathbf{S}}\hat{\mathbf{W}}_m)$, or $\mathbf{S}^{-1}\mathbf{s} \in \text{span}(\hat{\mathbf{W}}_m)$. Then, putting $\mathbf{s} = \mathbf{S}\hat{\mathbf{W}}_m\mathbf{q}$ in (7) for some \mathbf{q} , gives $\hat{\beta} = \hat{\mathbf{W}}_m\mathbf{q} = \mathbf{S}^{-1}\mathbf{s}$. □

5 The Bayes estimator

In Helland et al. (2012) a Bayes estimator under the PLS model was developed. The estimation was performed by a Markov Chain Monte Carlo approach. Specifically, for given m , and for observed centered data \mathbf{y} and \mathbf{X} the likelihood function is proportional to

$$f(\mathbf{y}, \mathbf{X} | \boldsymbol{\nu}, \boldsymbol{\gamma}, \mathbf{D}, \sigma^2) \propto (\sigma^2)^{-n/2} \exp \left(-\frac{1}{2\sigma^2} (\mathbf{y} - \mathbf{X} \sum_{i=1}^m \gamma_i \mathbf{d}_i)' (\mathbf{y} - \mathbf{X} \sum_{i=1}^m \gamma_i \mathbf{d}_i) \right) \times \left(\prod_{i=1}^p \nu_i \right)^{-n/2} \prod_{j=1}^n \exp \left(-\frac{1}{2} \mathbf{x}_j' \left(\sum_{i=1}^p \frac{1}{\nu_i} \mathbf{d}_i \mathbf{d}_i' \right) \mathbf{x}_j \right), \quad (8)$$

where $\boldsymbol{\nu} = [\nu_1, \dots, \nu_p]$ and $\mathbf{D} = [\mathbf{d}_1, \dots, \mathbf{d}_p]$ are the eigenvalues and the eigenvectors of the \mathbf{x} -covariance matrix $\boldsymbol{\Sigma}_{xx}$, and where $\boldsymbol{\gamma} = [\gamma_1, \dots, \gamma_m]$ are regression parameters of the PLS-model.

As argued in Helland et al. (2012), a near optimal equivariant regressor is found as the Bayesian estimator under rotation invariant prior for $\mathbf{d}_1, \dots, \mathbf{d}_p$ and prior $\pi(\boldsymbol{\gamma}) = \prod_i 1/\gamma_i^{1-\epsilon}$, where $1/\epsilon$ is a large uneven integer. Slightly modified scale priors are also chosen for $\boldsymbol{\nu}$ as $\pi(\boldsymbol{\nu}) = \prod_i 1/\nu_i \exp(-\epsilon_\nu/2\nu_i)$ and for σ^2 as $\pi(\sigma^2) = 1/\sigma^2 \exp(-\epsilon_\sigma/2\sigma^2)$. Here ϵ_ν and ϵ_σ are some small numbers chosen to ensure properness of the posterior distribution.

The R-code for Bayes PLS is available at <http://www.github.com/solvsa/BayesPLS>.

6 Data simulations for comparison of estimators

A comparative study of the prediction performances of the regular PLS algorithm, the maximum likelihood envelope method, the Bayes PLS and the method of ordinary least squares (OLS) was performed on data simulated from the random regression model (1). In this study, we consider envelope method for predictor reduction (Cook et al., 2013) and use R-code discussed in Cook et al. (2016). A detailed description of the simulation procedure can be found in (Sæbø et al., 2015) with the accompanying R-package `simrel`, but key features of the approach are presented next. The simulation set up is best explained from re-expressing model (1) as

$$\begin{bmatrix} y \\ \mathbf{x} \end{bmatrix} \sim \mathcal{N}(\boldsymbol{\mu}_{yx}, \boldsymbol{\Sigma}_{yx}) = \mathcal{N} \left(\begin{bmatrix} \mu_y \\ \boldsymbol{\mu}_x \end{bmatrix}, \begin{bmatrix} \sigma_y^2 & \boldsymbol{\sigma}_{xy}^t \\ \boldsymbol{\sigma}_{xy} & \boldsymbol{\Sigma}_{xx} \end{bmatrix} \right) \quad (9)$$

where, $\boldsymbol{\sigma}_{xy}$ is a vector holding the covariances between the predictors (\mathbf{x}) and the response

(y). The vector of regression coefficients β is by standard theory given as $\beta = \Sigma_{xx}^{-1} \sigma_{yx}$, which in turn can be expressed in terms of the eigenvalues ν_1, \dots, ν_p and the eigenvectors d_1, \dots, d_p of Σ_{xx} :

$$\beta = \sum_{i=1}^p \frac{d_i' \sigma_{yx}}{\nu_i} d_i = \sum_{i=1}^p \gamma_i d_i \quad (10)$$

as given in eq. (2). In `simrel` the following simplifying assumptions are made:

- ▷ It is assumed that $\nu_i = e^{-\eta(i-1)}$ for $i = 1, \dots, p$, implying $\nu_1 = 1$ (which we may assume without loss of generality,) and that all subsequent eigenvalues are decaying according to the size of the parameter η . A large η gives a rapid decrease in eigenvalues, implying high level of multi-collinearity in x .
- ▷ It is assumed that $m \leq p$ eigenvectors are relevant for y , which means that equation (2) (potentially) reduces to

$$\beta = \sum_{i \in \mathcal{P}} \gamma_i d_i \quad (11)$$

where m -vector \mathcal{P} is the set of indices of the relevant components for which $\gamma_i = d_i' \sigma_{yx} / \nu_i \neq 0$. Hence, the envelope or the relevant space has dimension m (see Theorem 1 above).

- ▷ Without loss of generality it is further assumed that $\sigma_y = 1$ and $\mu_{yx} = \mathbf{0}$.

In `simrel` the actual values of σ_{xy} were set to satisfy a pre-specified value of the population coefficient of determination ρ^2 . It may be shown that under the above assumptions $\rho^2 = \sigma_{xy}' \Sigma_{xx}^{-1} \sigma_{xy}$. This completes the specification of the parameters used in `simrel`, and in the present comparison study a design for the simulated data sets in terms of these parameters were as defined in Table 1 below.

From the possible combination of the above parameters, 32 calibration sets were simulated with 5 replications of each, i.e. there were 160 calibration sets (datasets) altogether.

Table-1 here

7 Systematic comparisons

A systematic comparison of the methods across the simulation designs was made on the basis of their ability to predict test samples. Since the distribution of the simulated variables is fully

known, the expected mean squared error of prediction (MSEP) based on some $\hat{\beta}$ estimated from a calibration set, may be found as:

$$E_x \left[E_y (y - \hat{y})^2 \right] = \left[\sigma^2 + E (\hat{\beta} - \beta)^t \Sigma_{xx} (\hat{\beta} - \beta) \right] \frac{n+1}{n}$$

in the model. The expectation on the right-hand side of the above expression is estimated for each method and for each design as an average over the five replicated calibration sets. In order to study the effects of p , ρ^2 , $\text{relpos}(\mathcal{P})$, Method, and (η) along with their interactions, we first retrieved the minimum MSEP for each method across 1 to 10 components (assumed sizes of the relevant space). In Figure 1 interaction plots for these data properties are displayed.

Figure-1 Here

The effect of the third order interaction between p , ρ^2 and Methods, which we see in Figure 1 (left), shows that the maximum likelihood based estimation methods, in our case, the envelope and the ordinary least squares, perform poorly on data sets with large number of variables and low ρ^2 . Still, the performance of the envelope is better than ordinary least squares also in situations where $p \sim n$. The interaction plots suggest that the Bayes PLS and ordinary PLS estimation methods are better and more stable on average than the two other methods.

Similarly, the effect of third order interaction between relpos , η and Method in Figure 1 (right) shows that OLS method gives higher prediction error than other methods, but the effect of relpos is small but notable for the envelope method. Again Bayes PLS and ordinary PLS are best.

The prediction error plots below are organized into four groups: a) $p = 15$, $\rho^2 = 0.5$; b) $p = 15$, $\rho^2 = 0.9$; c) $p = 40$, $\rho^2 = 0.5$ and d) $p = 40$, $\rho^2 = 0.9$. The ordinary least squares prediction error is shown by a straight dotted line.

Figure-2 Here

In group a), with small number of variables ($p \ll n$) and noisy data ($\rho^2 = 0.5$), Figure 2 shows that all the estimation methods performed better than ordinary least squares for all designs in this group, Bayes PLS being best in nearly all cases. Some convergence problems with Bayes PLS when eigenvalues decrease rapidly can be ignored since the minimum MSEP is already obtained from fewer components.

Figure-3 Here

Having few variables rich with information ($\rho^2 = 0.9$), the designs in group b) (Figure 3) leads to easy prediction with low prediction error in general for all methods. All the methods including OLS have small MSEs, but the other methods are still dominant. In most of the situations, Bayes PLS has reached minimum error with only one component. In this group, the performance of envelope is better than regular PLS and the minimum error for envelope is also achieved with fewer components.

Figure-4 Here

Low information content combined with many predictor variables characterize the designs in group c) and prediction is in general difficult for these designs. In Figure 4, the methods based on maximum likelihood estimation performed poorly and often poorer than an average guess. Bayes PLS and regular PLS performed well, as in the previous designs.

Figure-5 Here

Although having 40 predictors ($p \sim n$), all methods are performing well on the designs in group d) due to rich information (high ρ^2). Figure 5 shows that in most of the situations (except in design-16), the envelope method has nearly attained true minimum error (0.1) and has outperformed OLS. However, its prediction error is still larger than Bayes PLS and PLS. Bayes PLS and PLS methods are highly stable and are closer to true minimum error. Both of these methods have outperformed OLS and the envelope methods. Further, Bayes PLS is able to obtain its minimum prediction error with only few components.

In general ordinary PLS is very stable in all situations. It is extensible (lots of variants has been developed after its introduction), easy and less time consuming to fit than Bayes PLS and the envelope method. If the issue is to get closer prediction from squeezing information as much as possible, Bayes PLS will be an alternative. Its performance with fewer number of components is stable and better in all designs studied here. The envelope method performed better than OLS and the performance increased for informative data ($\rho^2 = 0.9$). However, the increased error with additional component indicates low shrinkage effect of the method.

Correlation between estimated and true regression coefficients (β) along with the mean square error of estimation (MSE) is presented in Figure 6. In case of ordinary PLS and the envelope method, the correlation for design 1 from group a) and design 3 from group b), both

having 15 predictors, is high for small components. However for design 2 from group c) and design 4 from group d), envelope methods exhibit sudden decrease in the correlation with corresponding increase in estimation error. The impressive prediction performance of Bayes PLS is also seen from the high correlation of estimated coefficients and true coefficients. In addition, the average MSE of regression for this methods is also small compared to others for all the components.

Although having low prediction error in case of envelope estimation method, the coefficient estimates are highly unstable for different components which we can see from its variation in correlation with true coefficients (Figure 6, top. Bayes PLS and regular PLS estimates are more stable over different replicates and for different components (Figure 6, bottom) especially when $p \sim n$. This stability agrees with the low prediction error we have discussed before.

Figure-6 Here

8 Discussion

There is a vast literature on applications of PLS, not only in chemistry, but in a large number of applied fields; see for instance Boulesteix and Strimmer (2007). Many further references are given in Mehmood and Ahmed (2016).

Sometimes the issue is prediction, but very often one also see interpretations of scoring plots, loading plots and correlation plots; see for instance Martens and Martens (2001). Such plots are not unfamiliar to statisticians in principal component connections, but they are much more used by the chemometric society and many scientists find them informative. They are plots of the sample variants of the latent variables and parameters defined by (3), (4) and (5), and thus involve consistent estimates of these quantities when $n \rightarrow \infty$ and probably also in the more general case $p/n \rightarrow 0$.

By comparison there are relatively few papers by mathematical statisticians investigating statistical properties of the partial least squares regression method itself. There are however several investigations on the shrinkage properties of PLS; see Krämer (2007) and references there, and also Foschi (2015) with references. Garthwaite (1994) offered a simple interpretation of PLS. Stone and Brooks (1990) and Naik and Tsai (2000) discuss different generalizations of PLS; in the latter paper also consistency of PLS is proved. Stoica and Söderström (1998) derives

asymptotic formulae related to PLS. Chun and Keleş (2010) extends consistency to the case $p/n \rightarrow 0$, introduces a sparse PLS algorithm, and compares methods by simulation. Krämer and Sugiyama (2012) discusses the degrees of freedom of PLS regression, and uses this concept in model selection. See also references in this last paper.

The purpose of the present article has been to discuss the approach to PLS-regression via model reduction in the random x multiple regression model, and to compare estimators in this reduced model.

From simulations, the Bayes estimator under the PLS model seems to have very good properties. In virtually all of the 32 designs, the MSEP curve for Bayes PLS lie below that for ordinary PLS and also that for the maximum likelihood envelope model. A particularly desirable feature of Bayes-PLS is that the MSEP-curve seems to be almost flat for small values of m . Thus the error made by choosing a wrong number of components m by crossvalidation must be expected to be small if m is relatively small.

Envelope and Bayes PLS estimation methods, when compared with ordinary PLS methods, display better prediction performance (only when p is small for the envelope method). However both of them have their disadvantages. The envelope method, as based on maximum likelihood, breaks down when p approaches n , while Bayes PLS has time consuming computation, and in our simulations it failed to converge for some cases.

For practical purposes the ordinary PLS algorithm still seems to be a good option.

References

- Boulesteix, A.-L. and Strimmer, K. (2007). Partial least squares: a versatile tool for the analysis of high-dimensional genomic data. *Briefings in bioinformatics*, 8(1):32–44.
- Chun, H. and Keleş, S. (2010). Sparse partial least squares regression for simultaneous dimension reduction and variable selection. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, 72(1):3–25.
- Chung, D. and Keles, S. (2010). Sparse partial least squares classification for high dimensional data. *Statistical applications in genetics and molecular biology*, 9(1).
- Cook, D., Su, Z., Yang, Y., et al. (2015). envlp: A matlab toolbox for computing envelope estimators in multivariate analysis. *Journal of Statistical Software*, 62(1):1–20.

- Cook, R., Helland, I., and Su, Z. (2013). Envelopes and partial least squares regression. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, 75(5):851–877.
- Cook, R. D., Forzani, L., and Su, Z. (2016). A note on fast envelope estimation. *Journal of Multivariate Analysis*, 150:42–54.
- Cook, R. D., Li, B., and Chiaromonte, F. (2010). Envelope models for parsimonious and efficient multivariate linear regression. *Statistica Sinica*, pages 927–960.
- Cook, R. D. and Zhang, X. (2015a). Foundations for envelope models and methods. *Journal of the American Statistical Association*, 110(510):599–611.
- Cook, R. D. and Zhang, X. (2015b). Simultaneous envelopes for multivariate linear regression. *Technometrics*, 57(1):11–25.
- Cook, R. D. and Zhang, X. (2016). Algorithms for envelope estimation. *Journal of Computational and Graphical Statistics*, 25(1):284–300.
- Foschi, P. (2015). The geometry of pls shrinkages. Technical report, University of Bologna.
- Frank, L. E. and Friedman, J. H. (1993). A statistical view of some chemometrics regression tools. *Technometrics*, 35(2):109–135.
- Garthwaite, P. H. (1994). An interpretation of partial least squares. *Journal of the American Statistical Association*, 89(425):122–127.
- Hastie, T., Tibshirani, R., and Friedman, J. (2009). The elements of statistical learning 2nd edition.
- Helland, I. S. (1988). On the structure of partial least squares regression. *Communications in statistics-Simulation and Computation*, 17(2):581–607.
- Helland, I. S. (1990). Partial least squares regression and statistical models. *Scandinavian Journal of Statistics*, pages 97–114.
- Helland, I. S. (2000). Model reduction for prediction in regression models. *Scandinavian journal of statistics*, 27(1):1–20.
- Helland, I. S. (2001). Reduction of regression models under symmetry. *Contemporary Mathematics*, 287:139–154.

- Helland, I. S. and Almøy, T. (1994). Comparison of prediction methods when only a few components are relevant. *Journal of the American Statistical Association*, 89(426):583–591.
- Helland, I. S., Saebø, S., Tjelmeland, H., et al. (2012). Near optimal prediction from relevant components. *Scandinavian Journal of Statistics*, 39(4):695–713.
- Krämer, N. (2007). An overview on the shrinkage properties of partial least squares regression. *Computational Statistics*, 22(2):249–273.
- Krämer, N. and Sugiyama, M. (2012). The degrees of freedom of partial least squares regression. *Journal of the American Statistical Association*.
- Martens, H. and Martens, M. (2001). Multivariate analysis of quality. an introduction.
- Martens, H. and Naes, T. (1992). *Multivariate calibration*. John Wiley & Sons.
- Mehmood, T. and Ahmed, B. (2016). The diversity in the applications of partial least squares: an overview. *Journal of Chemometrics*, 30(1):4–17.
- Munck, L., Jespersen, B. M., Rinnan, Å., Seefeldt, H., Engelsen, M. M., Nørgaard, L., and Engelsen, S. B. (2010). A physiochemical theory on the applicability of soft mathematical models—experimentally interpreted. *Journal of Chemometrics*, 24(7-8):481–495.
- Naes, T. and Helland, I. S. (1993). Relevant components in regression. *Scandinavian journal of statistics*, pages 239–250.
- Naik, P. and Tsai, C.-L. (2000). Partial least squares estimator for single-index models. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, 62(4):763–771.
- Sæbø, S., Almøy, T., and Helland, I. S. (2015). simrel—a versatile tool for linear model data simulation based on the concept of a relevant subspace and relevant predictors. *Chemometrics and Intelligent Laboratory Systems*, 146:128–135.
- Stoica, P. and Söderström, T. (1998). Partial least squares: A first-order analysis. *Scandinavian Journal of Statistics*, 25(1):17–24.
- Stone, M. and Brooks, R. J. (1990). Continuum regression: cross-validated sequentially constructed prediction embracing ordinary least squares, partial least squares and principal components regression. *Journal of the Royal Statistical Society. Series B (Methodological)*, pages 237–269.

Sundberg, R. (1999). Multivariate calibration—direct and indirect regression methodology. *Scandinavian Journal of Statistics*, 26(2):161–207.

Wold, S., Ruhe, A., Wold, H., and Dunn, III, W. (1984). The collinearity problem in linear regression. the partial least squares (pls) approach to generalized inverses. *SIAM Journal on Scientific and Statistical Computing*, 5(3):735–743.

Inge S. Helland

Department of Mathematics

University of Oslo

POBox 1053

NO-0316 Oslo

Norway

E-mail: ingeh@math.uio.no